

UROP: Evaluating SnapFuzz with ProFuzzBench

adotinthevoid.github.io/talks/snapfuzz.pdf

Alona Enraght-Moony

2022-09-09

Snapfuzz [1]

- ▶ Greybox Fuzzer for statefull network application
- ▶ Fork of AFLNet [4], itself a fork of AFL [5]
- ▶ Uses SaBRe [2] to intercept syscalls to dramatically increase speed.
 - ▶ Avoid synchronization delays
 - ▶ In memory filesystem
 - ▶ Optimize forkserver
 - ▶ Avoiding Sleep

ProFuzzBench [3]


- ▶ Benchmark suite for statefull fuzzers.
- ▶ 10 Protocols, 13 Implementations.
- ▶ Patches for Derandomization, Initial seeds, etc.
- ▶ Fuzzing effectiveness measured by coverage.

Various Problems Starting Running

- ▶ Building snapfuzz
- ▶ Memory Limit
- ▶ AFL ar
- ▶ Finding sabre
- ▶ Ubuntu 20.04 Glibc Debug symbols. ¹
- ▶ Profuzzbench Graphing Hard Coded
- ▶ Typoing 50 to 50/ silently gives wrong output
- ▶ Analysis/plotting script hardcoded to existing fuzzers

¹<https://bugs.launchpad.net/ubuntu/+source/glibc/+bug/1918035>

Perf Problems



Fuzzer	Execs	Execs/Sec	Edge Cov %	Line Cov %
aflnet	47139.25	6.55	40.35	59.35
snapfuzz	12765.75	1.77	37.60	55.7

Wrong Trees Barked Up

- ▶ CPU Cores
- ▶ AFL Flags

18.04 → 20.04 Glibc changes

Refactor nanosleep in terms of clock_nanosleep

author Adhemerval Zanella <adhemerval.zanella@linaro.org>
Tue, 5 Nov 2019 21:37:44 +0000 (21:37 +0000)
committer Adhemerval Zanella <adhemerval.zanella@linaro.org>
Wed, 6 Nov 2019 17:47:02 +0000 (14:47 -0300)

The generic version is straightforward. For Hurd, its nanosleep implementation is moved to clock_nanosleep with adjustments from generic unix implementation.

The generic clock_nanosleep unix version is also removed since it calls nanosleep.

Checked on x86_64-linux-gnu and powerpc64le-linux-gnu.

Reviewed-by: Florian Weimer <fweimer@redhat.com>

```
include/time.h
posix/nanosleep.c
sysdeps/mach/clock_nanosleep.c      [moved from sysdeps/unix/clock_nanosleep.c with 64% similarity]
sysdeps/mach/nanosleep.c            [deleted file]
sysdeps/unix/sysv/linux/clock_nanosleep.c
sysdeps/unix/sysv/linux/nanosleep.c [deleted file]
time/clock_nanosleep.c
```

3537ecb49cf7177274607004c562d6f9ecc99474


Fix

```
▼ 8 snapfuzz/main.c
```

...	@@ -811,6 +811,11 @@ int inanosleep(const struct timespec *req, struct timespec *rem) {
811	nanosleep((const struct timespec[]){0, 1L}, NULL);
812	return 0;
813	}
814	+ int iclock_nanosleep(clockid_t clockid, int flags,
815	+ const struct timespec *request, struct timespec *remain) {
816	+ clock_nanosleep(CLOCK_REALTIME, 0, (const struct timespec[]){0, 1L}, NULL);
817	+ return 0;
818	+ }
814	#endif // SF_SLEEP
815	
816	// static int cpus[8] = {0};
.... ↑	@@ -968,6 +973,9 @@ long handle_syscall(long sc_no, long arg1, long arg2, long arg3, long arg4,
968	#ifdef SF_SLEEP
969	} else if (sc_no == SYS_nanosleep) {
970	return inanosleep((const struct timespec *)arg1, (struct timespec *)arg2);
976	+ } else if (sc_no == SYS_clock_nanosleep) {
977	+ return iclock_nanosleep(arg1, arg2, (const struct timespec *)arg3,
978	+ (struct timespec *)arg4);
971	#endif // SF_SLEEP
972	// } else if (sc_no == SYS_getpid) {
973	// assert(false);
.... ↓	

ef005157cd97c9d3242b6d0a17908165fe1b74a6


Perf once fixed



Dnsmasq

Code coverage analysis for dnsmasq

Edge coverage over time (#edges)




Dnsmasq

fuzzer	run_no	time_spent	total_execs	ave_execs_per_sec	b_cov_percent	l_cov_percent
aflnet	1	28796.00	310781.00	10.79	14.60	23.70
aflnet	2	28796.00	322615.00	11.20	11.00	20.40
aflnet	3	28796.00	274408.00	9.53	5.80	11.60
aflnet	4	28796.00	430015.00	14.93	17.30	26.00
snapfuzz	1	28796.00	4493809.00	156.06	5.60	11.20
snapfuzz	2	28796.00	4292986.00	149.08	13.90	23.20
snapfuzz	3	28796.00	4955180.00	172.08	15.10	24.20
snapfuzz	4	28796.00	5248259.00	182.26	16.80	25.40
aflnet	average	28796.00	334454.75	11.61	12.18	20.42
snapfuzz	average	28796.00	4747558.50	164.87	12.85	21.00

LightFTP

Code coverage analysis for lightftp
Edge coverage over time (#edges)




LightFTP

fuzzer	run_no	time_spent	total_execs	ave_execs_per_sec	b_cov_percent	l_cov_percent
aflnet	1	28796.00	135028.00	4.69	43.10	63.00
aflnet	2	28796.00	137050.00	4.76	42.70	62.90
aflnet	3	28796.00	133225.00	4.63	43.20	63.30
aflnet	4	28796.00	163430.00	5.68	44.00	63.70
snapfuzz	1	28796.00	366503.00	12.73	45.10	64.50
snapfuzz	2	28796.00	150899.00	5.24	43.70	63.70
snapfuzz	3	28796.00	197958.00	6.87	43.70	63.40
snapfuzz	4	28796.00	142240.00	4.94	44.00	63.80
aflnet	average	28796.00	142183.25	4.94	43.25	63.22
snapfuzz	average	28796.00	214400.00	7.45	44.12	63.85

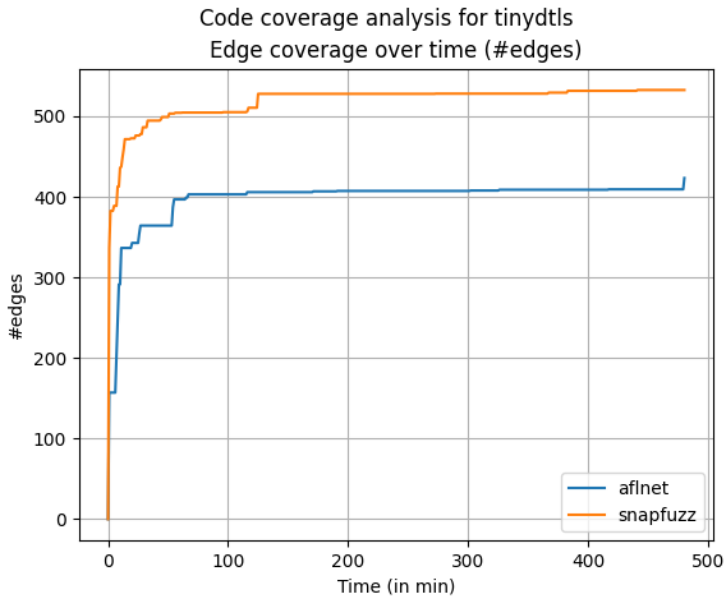
LIVE555

Code coverage analysis for live555
Edge coverage over time (#edges)



LIVE555

fuzzer	run_no	time_spent	total_execs	ave_execs_per_sec	b_cov_percent	l_cov_percent
aflnet	1	28796.00	302426.00	10.50	16.40	25.40
aflnet	2	28795.00	320816.00	11.14	16.60	25.90
aflnet	3	28796.00	332498.00	11.55	16.70	26.00
aflnet	4	28796.00	381030.00	13.23	16.70	26.00
snapfuzz	1	28796.00	1073514.00	37.28	16.90	26.10
snapfuzz	2	28796.00	1272257.00	44.18	17.20	26.30
snapfuzz	3	28795.00	1111905.00	38.61	16.70	25.50
snapfuzz	4	28796.00	1051298.00	36.51	16.70	25.90
aflnet	average	28795.75	334192.50	11.61	16.60	25.82
snapfuzz	average	28795.75	1127243.50	39.15	16.88	25.95



fuzzer	run_no	time_spent	total_execs	ave_execs_per_sec	b_cov_percent	l_cov_percent
aflnet	1	28795.00	53836.00	1.87	20.20	27.50
aflnet	2	28796.00	37449.00	1.30	24.00	32.10
aflnet	3	25622.00	41224.00	1.61	20.00	27.40
aflnet	4	28795.00	41222.00	1.43	24.90	32.60
snapfuzz	1	28796.00	5445009.00	189.09	26.60	34.00
snapfuzz	2	28796.00	5836083.00	202.67	28.00	34.70
snapfuzz	3	28796.00	6390305.00	221.92	31.20	43.40
snapfuzz	4	28795.00	5847171.00	203.06	26.40	33.80
aflnet	average	28002.00	43432.75	1.55	22.27	29.90
snapfuzz	average	28795.75	5879642.00	204.18	28.05	36.47

One's that didn't work

- ▶ Bftpd: `dup2(2)` suport missing
- ▶ forked-daapd: `statfs(2)` support missing
- ▶ DCMTK: Segfaults in startup code
- ▶ Kamilio: Deadlock between afl-fuzz and forkserver
- ▶ Exim: `wait4(-1, ...)` and `close(-1, ...)`, so forkserver exits.
- ▶ OpenSSH: Invalid pid sent to afl-fuzz
- ▶ OpenSSL: `close(3)` gives `EBADFD`, then forkserver exits.

Conclusions

- ▶ Snapfuzz isn't a drop in replacement for AFLNet.
- ▶ Snapfuzz is increadably fragile.
- ▶ It's not enough to be slightly faster.

Future work

- ▶ Add more syscalls
- ▶ Increase Debugability
 - ▶ Get PID of forkserver, and attach gdb before it crashes
- ▶ Fix snapfuzz bugs.

Source Code

- ▶ <https://github.com/aDotInTheVoid/profuzzbench/tree/snapfuzz>
- ▶ <https://github.com/aDotInTheVoid/snapfuzz-omni>
- ▶ <https://github.com/aDotInTheVoid/pfb-analysis>

Bibliography

[1]

Anastasios Andronidis and Cristian Cadar. 2022. SnapFuzz: High-throughput fuzzing of network applications. In *Proceedings of the 31st ACM SIGSOFT international symposium on software testing and analysis* (ISSTA 2022), Association for Computing Machinery, New York, NY, USA, 340–351. DOI:<https://doi.org/10.1145/3533767.3534376>

[2]

Paul-Antoine Arras, Anastasios Andronidis, Luís Pina, Karolis Mituzas, Qianyi Shu, Daniel Grumberg, and Cristian Cadar. 2022. SaBRe: Load-time selective binary rewriting. *International Journal on Software Tools for Technology Transfer* 24, 2 (April 2022), 205–223. DOI:<https://doi.org/10.1007/s10009-021-00644-w>

[3]

Roberto Natella and Van-Thuan Pham. 2021. ProFuzzBench: A benchmark for stateful protocol fuzzing. In *Proceedings of the 30th ACM SIGSOFT international symposium on software testing and analysis* (ISSTA 2021), Association for Computing Machinery, New York, NY, USA, 662–665. DOI:<https://doi.org/10.1145/3460319.3469077>

[4]

Van-Thuan Pham, Marcel Böhme, and Abhik Roychoudhury. 2020. AFLNET: A greybox fuzzer for network protocols. In *2020 IEEE 13th international conference on software testing, validation and verification (ICST)*, 460–465. DOI:<https://doi.org/10.1109/ICST46399.2020.00062>

[5]

Michał Zalewski. American fuzzy lop. Retrieved from <https://lcamtuf.coredump.cx/afl/>